

# REDIS

Finalist NoSQL techtalks

# INTRODUCTION

- Wikipedia: *Redis is an open-source, networked, in-memory, key-value data store with optional durability. It is written in ANSI C. The development of Redis is sponsored by VMware.*
- Easy!

# KEY-VALUE STORE?

- No, 5 different data structures
  - Strings (typical key-value)
  - Hashes
  - Lists
  - Sets
  - Sorted Sets
- Yes, each data structure uses key and values

# KEY-VALUE STORE?

- Relational databases only have tables, generic solution
  - Adds complexity and impacts performance
- Redis is not a one-size-fits-all solution.
- Use Redis data structures specific to your problem
  - Similar to the way you code...

# SERVER / CLIENT

- Spin up a server using redis-server
- Connect to it with redis-cli
- Simple!

# DATABASES

- Server has multiple database, referenced by number.
- Switch using SELECT index

# DATA STRUCTURES

# STRINGS: KEY-VALUE

- SET key value
- GET key
- Also (and more...)
  - MSET, MGET
  - GETSET
  - SETNX

# STRINGS: BITS

- SETBIT key offset value
- 2.6+
- BITCOUNT key [start] [end]
- BITOP operation destkey key
- See interesting metrics use case at:  
<http://blog.getspool.com/2011/11/29/fast-easy-realtime-metrics-using-redis-bitmaps/>

# STRINGS: COUNTERS

- INCR key
- INCRBY key increment
- DECR key
- DECRBY key decrement

# HASH

- HSET key field value
- HGET key field
- Also
  - HKEYS
  - HGETALL
  - HINCRBY

# LIST

- LPUSH en LPOP
- RPUSH en RPOP
- LRANGE
- RPOPLPUSH en BRPOPLPUSH (wtf?)

# SET

- Set operations (!)
- SADD key member [member...]
- SMEMBERS key
- SISMEMBER key
- Also...
  - SDIFF, SINTER and SUNION

# SORTED SET

- ZADD key score member [score] [member]
- ZRANGE key start stop [WITHSCORES]
- ZRANGEBYSCORE key min max

# MORE

- Publish / Subscribe
- Expiration
- Transactions
- Finding keys

# PERFORMANCE

# WHY?

- Specialized data structures, for example:
  - ISMEMBER command has  $O(1)$  complexity
  - See command reference for complexity values of different commands:  
<http://redis.io/commands>

# WHY?

- In-memory
- Redis does not care about values (also means: no queries)

# HOW FAST?

- Bulk import demo...

IS THAT ALL?

**WE NEED MORE POWER!**

# NORMAL MODE

- Client: INCR X
- Server: 1
- Client: INCR X
- Server: 2
- Client: INCR X
- Server: 3

# PIPELINING!

- Client: INCR X
- Client: INCR X
- Client: INCR X
- Server: 1
- Server: 2
- Server: 3



# HOW FAST?

- Pipelining enhanced bulk import demo...

IS THAT ALL?

# REDIS-CLI --PIPE

- Mass insertion from command line.
- Commands are checked!



# MORE CAPACITY

# REDIS CLUSTER

- Currently work in progress.
- Ruby driver offers distributed Redis capabilities:
  - Add extra (master) servers for more capacity
  - Keys are distributed over servers
  - Client calculates where keys are stored
  - CPU :-(

# HOW FAST?

- Clustered bulk import demo... (no pipelining)

# PERSISTENCE

# WHAT'S THE CATCH?

- Data structures
- Speed
- In-memory
- There must be something wrong with persistence!

# MEMORY-BOUNDED

- Database size bound to memory size
- Virtual memory support deprecated

# READ IT

- Must read post by antirez on Redis persistence:  
<http://oldblog.antirez.com/post/redis-persistence-demystified.html>

# DURABILITY

- 1: Client sends a write command to database  
*Data is in clients memory*
- 2: Database receives the write  
*Data is in server memory*
- 3: The database calls the system call that writes the data on disk  
*Data is in kernel's buffer*
- 4: The OS transfers the write buffer to the disk controller  
*Data is in the disk cache*
- 5: The disk controller actually writes the data into a physical media.  
*Persisted!*

# DURABILITY

- 1: Client sends a write command to database  
*Data is in clients memory*
- 2: Database receives the write  
*Data is in server memory*
- 3: The database calls the system call that writes the data on disk  
*Data is in kernel's buffer*
- 4: The OS transfers the write buffer to the disk controller  
*Data is in the disk cache*
- 5: The disk controller actually writes the data into a physical media.  
*Persisted!*

# DURABILITY

- Data to disk
- We can control:
  - Data written to OS kernel using `write` call
  - Data committed to disk using `fsync` call

# DURABILITY

- Redis has two persistence options:
  - RDB
  - AOF

RDB

# RDB

- Snapshotting
- Specify at what intervals you want a snapshot to be taken.
- Not very durable
- Easy backups

# AOF

# APPEND-ONLY-FILE

- Much more durable, but slower (in default mode)
- Buffer flushed (`write`) after each command (when re-entering the event loop)
- Different `fsync` policies (from worst to best):
  - none (OS decides)
  - every second (default)
  - every command (slow, but best durability in any db system)
- Append only == No corruption

# MASTER-SLAVE

# MASTER-SLAVE

- Master can have multiple slaves
- Slaves can have slaves as well
- Replication is non-blocking on master as well as slave

# HANDS-ON

# TWITTER

- `git clone https://github.com/robinroestenburg/redis-techtalk`
- `cd sinatra-twitter`
- `bundle`
- `rackup`
- Go to <http://localhost:9292> and sign up/in.

# TWITTER

- Implement all methods that are now returning dummy or empty results.
- Extra points:
  - Implement favoriting a Tweet and viewing favorites.
  - When following a user, all his tweets must be visible in the followers' timeline (in the right order).